

# EAS508-HW4

Saish Mandavkar

2022-10-20

## Lab Code Homework

### 5.3 Cross Validation Labs

#### 5.3.1 Validation Set Approach

```
# Setting the seed and loading the data
```

```
library(ISLR2)
```

```
## Warning: package 'ISLR2' was built under R version 4.0.5
```

```
set.seed(1)
```

```
train <- sample(392,196)
```

```
# Fitting a linear regression on the train data using subset option
```

```
lm.fit <- lm(mpg ~ horsepower, data = Auto, subset = train)
```

```
# Predicting the estimates for the 392 observations and calculate the MSE for 192 observations
```

```
mean((Auto$mpg - predict(lm.fit, Auto))[-train]^2)
```

```
## [1] 23.26601
```

```
# Fitting cubic regression and calculating the MSE
```

```
lm.fit2 <- lm(mpg ~poly(horsepower, 2), data = Auto, subset = train)
```

```
mean((Auto$mpg - predict(lm.fit2, Auto))[-train]^2)
```

```
## [1] 18.71646
```

```
# Fitting quadratic regression and calculating the MSE
```

```
lm.fit3 <- lm(mpg ~ poly(horsepower, 3), data = Auto, subset = train)
```

```
mean((Auto$mpg - predict(lm.fit3, Auto))[-train]^2)
```

```
## [1] 18.79401
```

```
# Using different seed and calculatiing the values for all the three regressions - will result into dif
```

```
set.seed(2)
train <- sample(392,196)

# Linear regression MSE

lm.fit <- lm(mpg ~ horsepower, data = Auto, subset = train)

mean((Auto$mpg - predict(lm.fit, Auto))[-train]^2)
```

```
## [1] 25.72651
```

```
# Cubic regression MSE
```

```
lm.fit2 <- lm(mpg ~poly(horsepower, 2), data = Auto, subset = train)

mean((Auto$mpg - predict(lm.fit2, Auto))[-train]^2)
```

```
## [1] 20.43036
```

```
# Quadratic regression MSE
```

```
lm.fit3 <- lm(mpg ~poly(horsepower, 3), data = Auto, subset = train)

mean((Auto$mpg - predict(lm.fit3, Auto))[-train]^2)
```

```
## [1] 20.38533
```

### 5.3.2 Leave One-Out Cross-Validation

```
# LOOCV using glm() package
```

```
glm.fit <- glm(mpg ~ horsepower, data = Auto)

coef(glm.fit)
```

```
## (Intercept)  horsepower
## 39.9358610  -0.1578447
```

```
# LOOCV using normal lm() function
```

```
lm.fit <- lm(mpg ~ horsepower, data = Auto)

coef(lm.fit)
```

```
## (Intercept) horsepower
## 39.9358610 -0.1578447
```

```
# Cross-validation error using glm() package
```

```
library(boot)
```

```
## Warning: package 'boot' was built under R version 4.0.5
```

```
glm.fit <- glm(mpg ~ horsepower, data = Auto)
```

```
cv.err <- cv.glm(Auto, glm.fit)
```

```
cv.err$delta
```

```
## [1] 24.23151 24.23114
```

```
# Calculating CV error for polynomial of order 1 to 10 using a for loop.
```

```
cv.error <- rep(0,10)
```

```
for (i in 1:10) {
```

```
  glm.fit <- glm(mpg ~ poly(horsepower, i), data = Auto)
  cv.error[i] <- cv.glm(Auto, glm.fit)$delta[1]
```

```
}
```

```
cv.error
```

```
## [1] 24.23151 19.24821 19.33498 19.42443 19.03321 18.97864 18.83305 18.96115
## [9] 19.06863 19.49093
```

### 5.3.3 k-Fold Cross Validation

```
# Calculating k-fold CV error for polynomial of order 1 to 10 with k = 10
```

```
set.seed(17)
```

```
cv.error.10 <- rep(0,10)
```

```
for (i in 1:10) {
```

```
  glm.fit <- glm(mpg ~ poly(horsepower, i), data = Auto)
  cv.error.10[i] <- cv.glm(Auto, glm.fit, K = 10)$delta[1]
```

```
}
```

```
cv.error.10
```

```
## [1] 24.27207 19.26909 19.34805 19.29496 19.03198 18.89781 19.12061 19.14666
## [9] 18.87013 20.95520
```

### 6.5.3 PCR and PLS Regression

```
# Creating model matrix for x and storing all salary values in y

# Omitting NA values

Hitters <- na.omit(Hitters)

x <- model.matrix(Salary ~ ., Hitters)[, -1]
y <- Hitters$Salary

# Creating train and test by setting the R seed

set.seed(1)
train <- sample(1:nrow(x), nrow(x) / 2)
test <- (-train)
y.test <- y[test]
```

```
# Applying PCR to Hitters data to predict Salary

library(pls)
```

#### Principal Components Regression

```
##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
##   loadings
```

```
set.seed(2)
pcr.fit <- pcr(Salary ~., data = Hitters, scale = TRUE, validation = "CV")
```

```
# Checking summary of our fit


summary(pcr.fit)
```

```
## Data:      X dimension: 263 19
## Y dimension: 263 1
## Fit method: svdpc
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              452    351.9   353.2   355.0   352.8   348.4   343.6
## adjCV           452    351.6   352.7   354.4   352.1   347.6   342.7
```

```
##      7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV      345.5   347.7   349.6   351.4   352.1   353.5   358.2
## adjCV    344.7   346.7   348.5   350.1   350.7   352.0   356.5
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV      349.7   349.4   339.9   341.6   339.2   339.6
## adjCV    348.0   347.7   338.2   339.7   337.2   337.6
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X      38.31   60.16   70.84   79.03   84.29   88.63   92.26   94.96
## Salary  40.63   41.58   42.17   43.22   44.90   46.48   46.69   46.75
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X      96.28   97.26   97.98   98.65   99.15   99.47   99.75
## Salary  46.86   47.76   47.82   47.85   48.10   50.40   50.55
##      16 comps 17 comps 18 comps 19 comps
## X      99.89   99.97   99.99   100.00
## Salary  53.01   53.85   54.61   54.61
```

```
# Plotting cross-validation MSE
```

```
validationplot(pcr.fit, val.type = "MSEP")
```




```
# Performing PCR on the training data using subset function and plotting the CV MSE
```

```
set.seed(1)
```

```

pcr.fit <- pcr(Salary ~., data = Hitters, subset = train, scale = TRUE, validation = "CV")
validationplot(pcr.fit, val.type = "MSEP")

```



```

# Find the lowest CV error when M = 5

```

```

pcr.pred <- predict(pcr.fit, x[test, ], ncomp = 5)
mean((pcr.pred - y.test)^2)

```

```

## [1] 142811.8

```

```

# Fit PCR on complete dataset using M = 5 identified by CV

```

```

pcr.fit <- pcr(y~x, scale = TRUE, ncomp = 5)
summary(pcr.fit)

```

```

## Data:      X dimension: 263 19
## Y dimension: 263 1
## Fit method: svdpc
## Number of components considered: 5

```

```
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps
## X      38.31   60.16   70.84   79.03   84.29
## y      40.63   41.58   42.17   43.22   44.90
```

```
# Implement PLS using plsr() function
```

```
set.seed(1)
pls.fit <- plsr(Salary ~ ., data = Hitters, subset = train, scale = TRUE, validation = "CV")
summary(pls.fit)
```

## Partial Least Squares

```
## Data:      X dimension: 131 19
## Y dimension: 131 1
## Fit method: kernelpls
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           428.3    325.5    329.9    328.8    339.0    338.9    340.1
## adjCV         428.3    325.0    328.2    327.2    336.6    336.1    336.6
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           339.0    347.1    346.4    343.4    341.5    345.4    356.4
## adjCV         336.2    343.4    342.8    340.2    338.3    341.8    351.1
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV           348.4    349.1    350.0    344.2    344.5    345.0
## adjCV         344.2    345.0    345.9    340.4    340.6    341.1
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X           39.13   48.80   60.09   75.07   78.58   81.12   88.21   90.71
## Salary       46.36   50.72   52.23   53.03   54.07   54.77   55.05   55.66
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X           93.17   96.05   97.08   97.61   97.97   98.70   99.12
## Salary       55.95   56.12   56.47   56.68   57.37   57.76   58.08
##      16 comps 17 comps 18 comps 19 comps
## X           99.61   99.70   99.95   100.00
## Salary       58.17   58.49   58.56   58.62
```

```
# Evaluating the corresponding test set MSE
```

```
pls.pred <- predict(pls.fit, x[test, ], ncomp = 1)
mean((pls.pred - y.test)^2)
```

```
## [1] 151995.3
```

```
# Fitting PLS on complete dataset when M = 1

pls.fit <- plsr(Salary ~ ., data = Hitters, scale = TRUE, ncomp = 1)

summary(pls.fit)
```

```
## Data:      X dimension: 263 19
## Y dimension: 263 1
## Fit method: kernelpls
## Number of components considered: 1
## TRAINING: % variance explained
##           1 comps
## X           38.08
## Salary      43.05
```

### 7.8.3 GAMs

```
# Fit a GAM or predict wage using natural spline functions of years and age.

gam1 <- lm(wage ~ splines::ns(year, 4) + splines::ns(age, 5) + education, data = Wage)
```

```
# Fit the model using smoothing splines
```

```
library(gam)
```

```
## Warning: package 'gam' was built under R version 4.0.5
```

```
## Loading required package: splines
```

```
## Loading required package: foreach
```

```
## Warning: package 'foreach' was built under R version 4.0.5
```


```
## Loaded gam 1.20.1
```

```
gam.m3 <- gam(wage ~ s(year, 4) + s(age, 5) + education, data = Wage)
```

```
# Plot the model
```


```
par(mfrow = c(1,3))
plot(gam.m3, se = TRUE, col = "blue")
```





```
# Plotting the GAM created using lm

par(mfrow = c(1,3))
plot.Gam(gam1, se = TRUE, col = "red")
```



```
# Performing ANOVA test to determine the best model
```

```
gam.m1 <- gam(wage ~ s(age, 5) + education, data = Wage)
gam.m2 <- gam(wage ~ year + s(age, 5) + education, data = Wage)

anova(gam.m1, gam.m2, gam.m3, test = "F")
```

```
## Analysis of Deviance Table
```

```
##
```

```
## Model 1: wage ~ s(age, 5) + education
```

```
## Model 2: wage ~ year + s(age, 5) + education
```

```
## Model 3: wage ~ s(year, 4) + s(age, 5) + education
```

```
##   Resid. Df Resid. Dev Df Deviance      F    Pr(>F)
```

```
## 1      2990      3711731
```

```
## 2      2989      3693842  1  17889.2 14.4771 0.0001447 ***
```

```
## 3      2986      3689770  3   4071.1  1.0982 0.3485661
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# summary of gam.m3
```

```
summary(gam.m3)
```

```
##
```

```
## Call: gam(formula = wage ~ s(year, 4) + s(age, 5) + education, data = Wage)
```

```
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -119.43  -19.70   -3.33   14.17  213.48
##
## (Dispersion Parameter for gaussian family taken to be 1235.69)
##
##      Null Deviance: 5222086 on 2999 degrees of freedom
## Residual Deviance: 3689770 on 2986 degrees of freedom
## AIC: 29887.75
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##           Df Sum Sq Mean Sq F value    Pr(>F)
## s(year, 4)   1   27162    27162  21.981 2.877e-06 ***
## s(age, 5)    1  195338   195338 158.081 < 2.2e-16 ***
## education    4 1069726   267432  216.423 < 2.2e-16 ***
## Residuals 2986 3689770     1236
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##           Npar Df Npar F    Pr(F)
## (Intercept)
## s(year, 4)         3  1.086 0.3537
## s(age, 5)          4 32.380 <2e-16 ***
## education
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Using the predict method for class GAM
```


```
preds <- predict(gam.m2, newdata = Wage)
```

```
# Using local regression fits in GAM using lo()
```

```
gam.lo <- gam(wage ~ s(year, df = 4) + lo(age, span = 0.7) + education, data = Wage)
```

```
par(mfrow = c(1,3))
```

```
plot.Gam(gam.lo, se = TRUE, col = "green")
```



```
# Using lo() to create interactions before calling gam
```

```
gam.lo.i <- gam(wage ~ lo(year, age, span = 0.5) + education, data = Wage)
```

```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : liv
## too small. (Discovered by lowesd)
```

```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : lv
## too small. (Discovered by lowesd)
```

```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : liv
## too small. (Discovered by lowesd)
```


```
## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame, bf.maxit, : lv
## too small. (Discovered by lowesd)
```

```
# Plotting the 2D surface using akima package
```


```
library(akima)
```

```
## Warning: package 'akima' was built under R version 4.0.5
```

```
par(mfrow = c(1,2))
plot(gam.lo.i)
```




1. < HS Grad    4. College Grad



*# Fitting a logistic regression GAM using I() function*

```
gam.lr <- gam(I(wage > 250) ~ year + s(age, df = 5) + education, family = binomial, data = Wage)
par(mfrow = c(1,3))
plot(gam.lr, se = TRUE, col = "green")
```



```
# Create a table with high earners in the < HS category
```

```
attach(Wage)
table(education, I(wage > 250))
```

```
##
## education      FALSE TRUE
## 1. < HS Grad    268    0
## 2. HS Grad      966    5
## 3. Some College 643    7
## 4. College Grad 663   22
## 5. Advanced Degree 381  45
```

```
# Fitting a logsitic regression GAM by skipping the education category
```

```
gam.lr.s <- gam( I(wage > 250) ~ year + s(age, df = 5) + education, family = binomial, subset = (educat.
par(mfrow = c(1,3))
plot(gam.lr.s, se = TRUE, col = "green")
```

